

# Word Embeddings

**CS 5624: Natural Language Processing**  
*Spring 2025*

<https://tuvllms.github.io/nlp-spring-2025>

**Tu Vu**



# Logistics

-  Homework 0 & Quiz 0 due tomorrow 

# Deep Research

- <https://openai.com/index/introducing-deep-research/>

# OmniHuman-1

- Albert Einstein
  - <https://x.com/stillgray/status/1887111258245095508>
- Taylor Swift singing Naruto song convincingly
  - <https://x.com/bevenky/status/1886840149012607087/video/1>

# s1: Simple test-time scaling

## s1: Simple test-time scaling

Niklas Muennighoff<sup>\*134</sup> Zitong Yang<sup>\*1</sup> Weijia Shi<sup>\*2</sup> Xiang Lisa Li<sup>\*1</sup> Li Fei-Fei<sup>1</sup> Hannaneh Hajishirzi<sup>23</sup>  
Luke Zettlemoyer<sup>2</sup> Percy Liang<sup>1</sup> Emmanuel Candès<sup>1</sup> Tatsunori Hashimoto<sup>1</sup>

### Abstract

Test-time scaling is a promising new approach to language modeling that uses extra test-time compute to improve performance. Recently, OpenAI’s o1 model showed this capability but did not publicly share its methodology, leading to many replication efforts. We seek the simplest approach to achieve test-time scaling and strong reasoning performance. First, we curate a small dataset **s1K** of 1,000 questions paired with reasoning traces relying on three criteria we validate through ablations: difficulty, diversity, and quality. Second, we develop budget forcing to control test-time compute.




Figure 1. **Test-time scaling with s1-32B.** We benchmark s1-32B on reasoning-intensive tasks and vary test-time compute.

CLJ 3 Feb 2025

<https://arxiv.org/pdf/2501.19393>

# RLHF

- <https://rlhfbook.com/>

# Cosine similarity

$$\text{dot product}(\mathbf{v}, \mathbf{w}) = \mathbf{v} \cdot \mathbf{w} = \sum_{i=1}^N v_i w_i = v_1 w_1 + v_2 w_2 + \dots + v_N w_N$$

$$\text{cosine}(\mathbf{v}, \mathbf{w}) = \frac{\mathbf{v} \cdot \mathbf{w}}{|\mathbf{v}| |\mathbf{w}|} = \frac{\sum_{i=1}^N v_i w_i}{\sqrt{\sum_{i=1}^N v_i^2} \sqrt{\sum_{i=1}^N w_i^2}}$$

# Word2vec

- **Continuous bag-of-words (CBow)**
- **Skip-gram with negative sampling (SGNS)**



# Skip-gram with negative sampling (SGNS)

- Simplifies the task: *binary classification* instead of *word prediction*
- Simplifies the architecture: *logistic regression* instead of *multi-layer neural network*


# Skip-gram with negative sampling (SGNS) (cont'd)

- Intuition
  - Treat the target word and a neighboring context word as *positive examples*
  - Randomly sample other words in the lexicon to get *negative samples*
  - Train a classifier (**self-supervision**) to distinguish those two cases
  - Use the learned weights as the embeddings


# Logistic regression

$$P(+|w, c) = \sigma(\mathbf{c} \cdot \mathbf{w}) = \frac{1}{1 + \exp(-\mathbf{c} \cdot \mathbf{w})}$$

the probability that  
word  $c$  is a real context  
word for target word  $w$



# Sigmoid function



## Sigmoid function (cont'd)

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

For  $1 - \sigma(x)$ :

$$1 - \sigma(x) = \frac{e^{-x}}{1 + e^{-x}}$$

Dividing numerator and denominator by  $e^{-x}$ :

$$1 - \sigma(x) = \frac{1}{e^x + 1} = \sigma(-x)$$

## Logistic regression (cont'd)

$$P(+|w, c) = \sigma(\mathbf{c} \cdot \mathbf{w}) = \frac{1}{1 + \exp(-\mathbf{c} \cdot \mathbf{w})}$$

the probability that word  $c$  is a real context word for target word  $w$

$$\begin{aligned} P(-|w, c) &= 1 - P(+|w, c) \\ &= \sigma(-\mathbf{c} \cdot \mathbf{w}) = \frac{1}{1 + \exp(\mathbf{c} \cdot \mathbf{w})} \end{aligned}$$

the probability that word  $c$  is **not** a real context word for target word  $w$

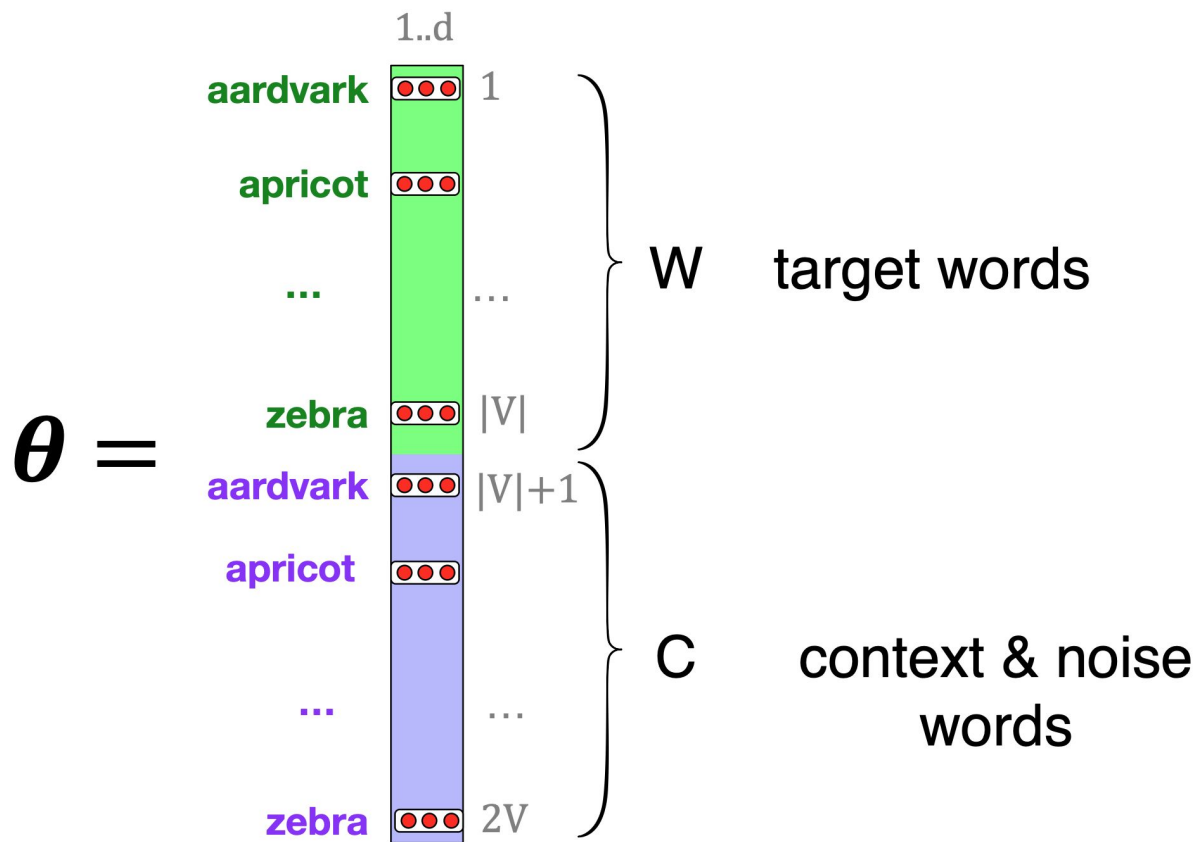
# The probability for many context words

Given a target word  $w$  and its context window of  $L$  words  $c_{1:L}$ , assigns a probability based on how similar this context window is to the target word

$$P(+|w, c_{1:L}) = \prod_{i=1}^L \sigma(\mathbf{c}_i \cdot \mathbf{w})$$

$$\log P(+|w, c_{1:L}) = \sum_{i=1}^L \log \sigma(\mathbf{c}_i \cdot \mathbf{w})$$

# Two embedding tables





# Learning skip-gram embeddings

... lemon, a [tablespoon of apricot jam, a] pinch ...  
c1 c2 w c3 c4

This example has a target word  $w$  (apricot), and 4 context words in the  $L = \pm 2$  window, resulting in 4 positive training instances (on the left below):

## positive examples +

$w$	$c_{\text{pos}}$
apricot	tablespoon
apricot	of
apricot	jam
apricot	a

## negative examples -

$w$	$c_{\text{neg}}$	$w$	$c_{\text{neg}}$
apricot	aardvark	apricot	seven
apricot	my	apricot	forever
apricot	where	apricot	dear
apricot	coaxial	apricot	if

# Loss function

Maximize the similarity of the target word, context word pairs  $(w, c_{\text{pos}})$  drawn from the positive examples

$$\begin{aligned} L &= -\log \left[ P(+|w, c_{\text{pos}}) \prod_{i=1}^k P(-|w, c_{\text{neg}_i}) \right] \\ &= - \left[ \log P(+|w, c_{\text{pos}}) + \sum_{i=1}^k \log P(-|w, c_{\text{neg}_i}) \right] \\ &= - \left[ \log P(+|w, c_{\text{pos}}) + \sum_{i=1}^k \log (1 - P(+|w, c_{\text{neg}_i})) \right] \\ &= - \left[ \log \sigma(c_{\text{pos}} \cdot w) + \sum_{i=1}^k \log \sigma(-c_{\text{neg}_i} \cdot w) \right] \end{aligned}$$

## Loss function (cont'd)

Minimize the similarity of the  $(w, c_{\text{neg}})$  pairs from the negative examples

$$\begin{aligned} L &= -\log \left[ P(+|w, c_{pos}) \prod_{i=1}^k P(-|w, c_{neg_i}) \right] \\ &= - \left[ \log P(+|w, c_{pos}) + \sum_{i=1}^k \log P(-|w, c_{neg_i}) \right] \\ &= - \left[ \log P(+|w, c_{pos}) + \sum_{i=1}^k \log (1 - P(+|w, c_{neg_i})) \right] \\ &= - \left[ \log \sigma(c_{pos} \cdot w) + \sum_{i=1}^k \log \sigma(-c_{neg_i} \cdot w) \right] \end{aligned}$$

# Training

$$\frac{\partial L}{\partial \mathbf{c}_{pos}} = [\sigma(\mathbf{c}_{pos} \cdot \mathbf{w}) - 1] \mathbf{w}$$

$$\frac{\partial L}{\partial \mathbf{c}_{neg}} = [\sigma(\mathbf{c}_{neg} \cdot \mathbf{w})] \mathbf{w}$$

$$\frac{\partial L}{\partial \mathbf{w}} = [\sigma(\mathbf{c}_{pos} \cdot \mathbf{w}) - 1] \mathbf{c}_{pos} + \sum_{i=1}^k [\sigma(\mathbf{c}_{neg_i} \cdot \mathbf{w})] \mathbf{c}_{neg_i}$$

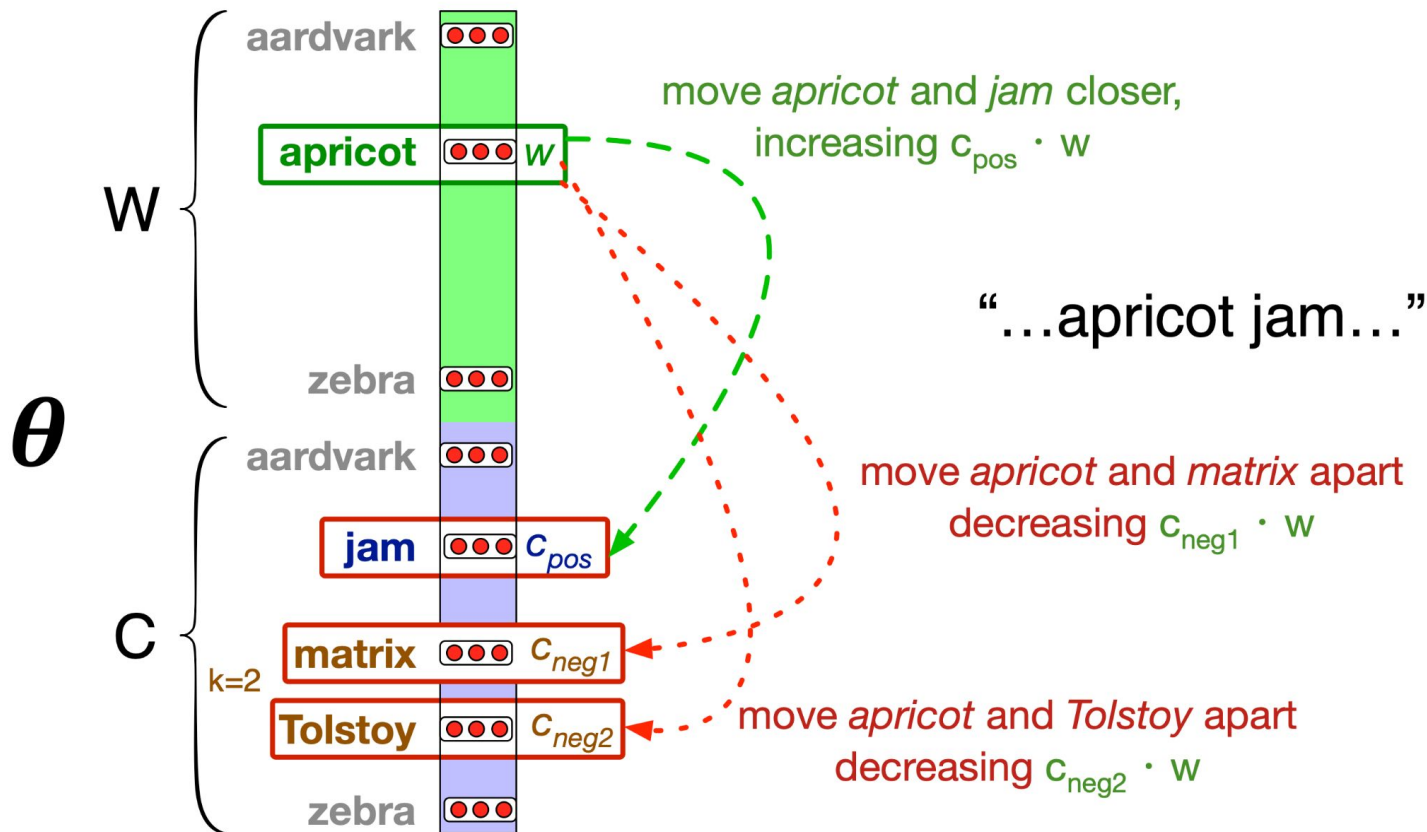
The update equations going from time step  $t$  to  $t + 1$  in stochastic gradient descent are thus:

$$\mathbf{c}_{pos}^{t+1} = \mathbf{c}_{pos}^t - \eta [\sigma(\mathbf{c}_{pos}^t \cdot \mathbf{w}^t) - 1] \mathbf{w}^t$$

$$\mathbf{c}_{neg}^{t+1} = \mathbf{c}_{neg}^t - \eta [\sigma(\mathbf{c}_{neg}^t \cdot \mathbf{w}^t)] \mathbf{w}^t$$

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \left[ [\sigma(\mathbf{c}_{pos}^t \cdot \mathbf{w}^t) - 1] \mathbf{c}_{pos}^t + \sum_{i=1}^k [\sigma(\mathbf{c}_{neg_i}^t \cdot \mathbf{w}^t)] \mathbf{c}_{neg_i}^t \right]$$

# Training (cont'd)




# Visualizing word embeddings



A two-dimensional (t-SNE) projection of embeddings for some words and phrases, showing that words with similar meanings are nearby in space.


# Semantic properties of word embeddings



$$\vec{\text{king}} - \vec{\text{man}} \approx \vec{\text{queen}} - \vec{\text{woman}}$$


$$\vec{\text{king}} - \vec{\text{man}} + \vec{\text{woman}} \approx \vec{\text{queen}}$$

# Semantic properties of word embeddings (cont'd)



$$\vec{\text{Paris}} - \vec{\text{France}} \approx \vec{\text{Rome}} - \vec{\text{Italy}}$$

$$\vec{\text{Paris}} - \vec{\text{France}} + \vec{\text{Italy}} \approx \vec{\text{Rome}}$$




$$\vec{\text{tree}} - \vec{\text{apple}} \approx \vec{\text{vine}} - \vec{\text{grape}}$$


$$\vec{\text{tree}} - \vec{\text{apple}} + \vec{\text{grape}} \approx \vec{\text{vine}}$$



# Semantic properties of word embeddings (cont'd)





# Semantic properties of word embeddings (cont'd)



***comparative  
& superlative  
morphology***


# Analogy problem $\mathbf{a} : \mathbf{b} :: \mathbf{a}^* : \mathbf{b}^*$

$a$  is to  $b$  as  $a^*$  is to what?



$$\hat{\mathbf{b}}^* = \operatorname{argmin}_{\mathbf{x}} \operatorname{distance}(\mathbf{x}, \mathbf{b} - \mathbf{a} + \mathbf{a}^*)$$



# Historical Semantics



A t-SNE visualization of the semantic change of 3 words in English using word2vec vectors.

# Biases in word embeddings

- Gender stereotypes



## Biases in word embeddings (cont'd)

- People in the US associate:
  - African-American names with unpleasant words (more than European-American names)
  - male names more with mathematics and female names with the arts, and old people's names with unpleasant words
- ...

# Word embedding methods

- Word2vec
- Glove: <https://nlp.stanford.edu/projects/glove/>
- Fasttext: <https://fasttext.cc/>

**Thank you!**