

# Scaling LLM Pretraining



**CS 5624: Natural Language Processing**  
*Spring 2025*

<https://tuvllms.github.io/nlp-spring-2025>

**Tu Vu**



# Logistics

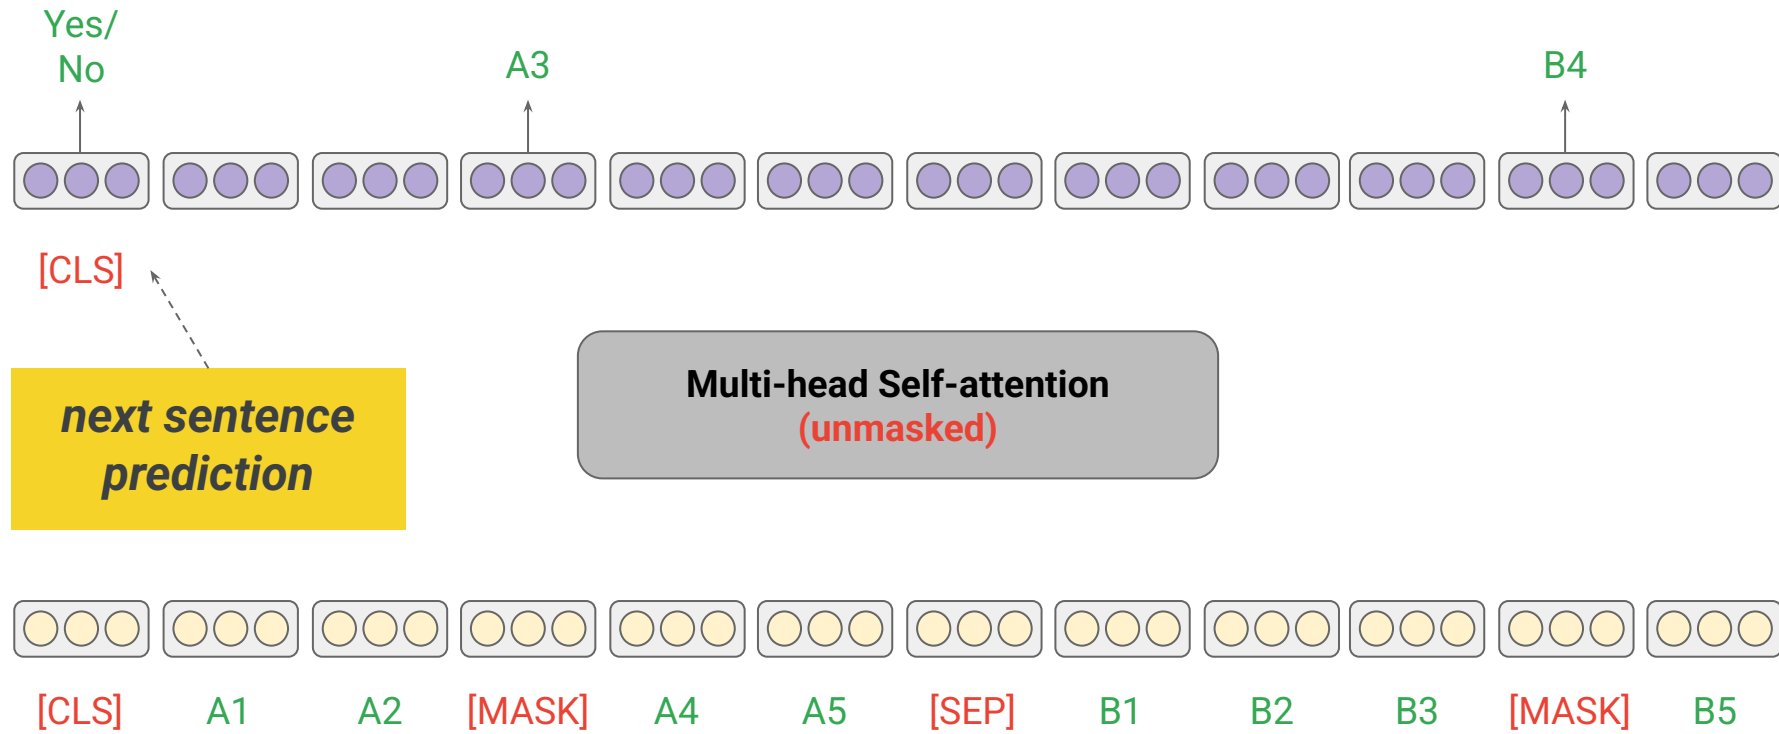
- Homework 1 & Quiz 1 will be released tomorrow
-  Final project proposal due on February 28 
  - Template is on Piazza

# BERT review

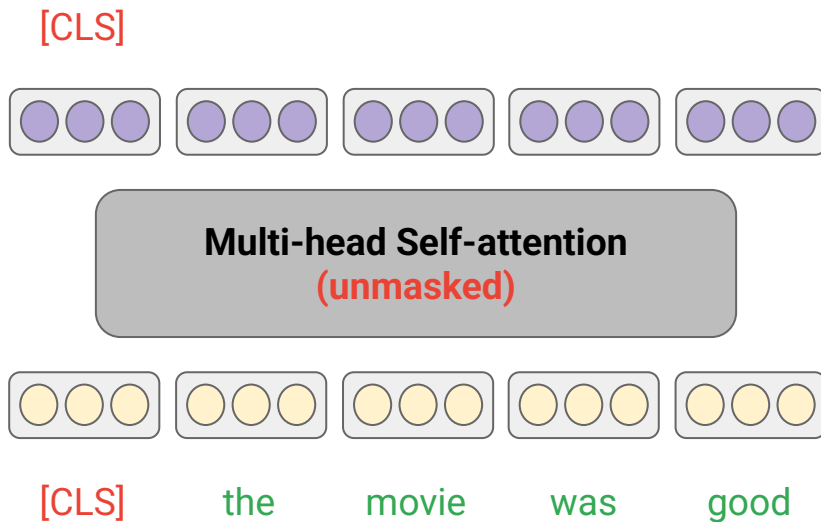

# Different model architectures

- Encoder-only
  - BERT
- Encoder-decoder
  - **T5**
- Decoder-only
  - GPT

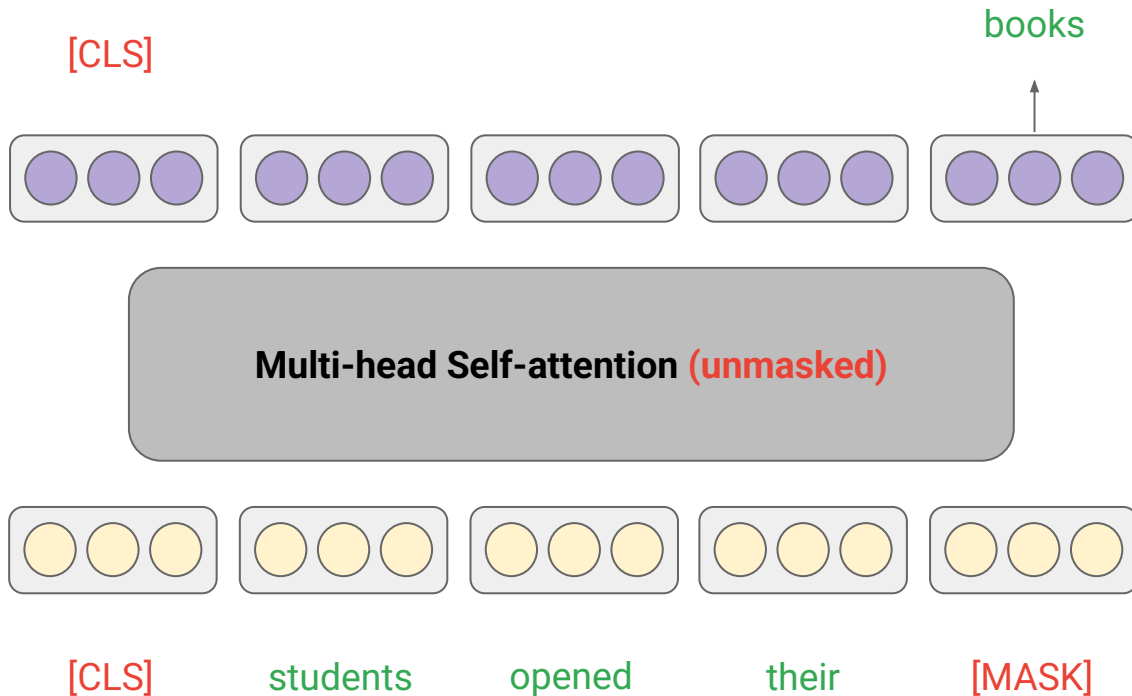
# BERT Pretraining



# BERT Fine-tuning



# Can BERT be used for text generation?



*iterative  
masking and  
unmasking*

# T5: Text-to-Text Transfer Transformer

## Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer

**Colin Raffel\***

CRAFFEL@GMAIL.COM

**Noam Shazeer\***

NOAM@GOOGLE.COM

**Adam Roberts\***

ADAROB@GOOGLE.COM

**Katherine Lee\***

KATHERINELEE@GOOGLE.COM

**Sharan Narang**

SHARANNARANG@GOOGLE.COM

**Michael Matena**

MMATENA@GOOGLE.COM

**Yanqi Zhou**

YANQIZ@GOOGLE.COM

**Wei Li**

MWEILI@GOOGLE.COM

**Peter J. Liu**

PETERJLIU@GOOGLE.COM

*Google, Mountain View, CA 94043, USA*



# T5 Pretraining: Span corruption

**<X>, <Y>: sentinel tokens**




**Transformer encoder**  
(unmasked)

Thank you <X> me to your party <Y> week

*encoder*

<X> for inviting <Y> last <EOS>



<BOS> <X> for inviting <Y> last

*decoder*

Thank you ~~for inviting~~ me to your party ~~last~~ week

# T5 Fine-tuning



**Transformer encoder**  
(unmasked)

sentiment analysis: this movie was good

*encoder*

positive <EOS>

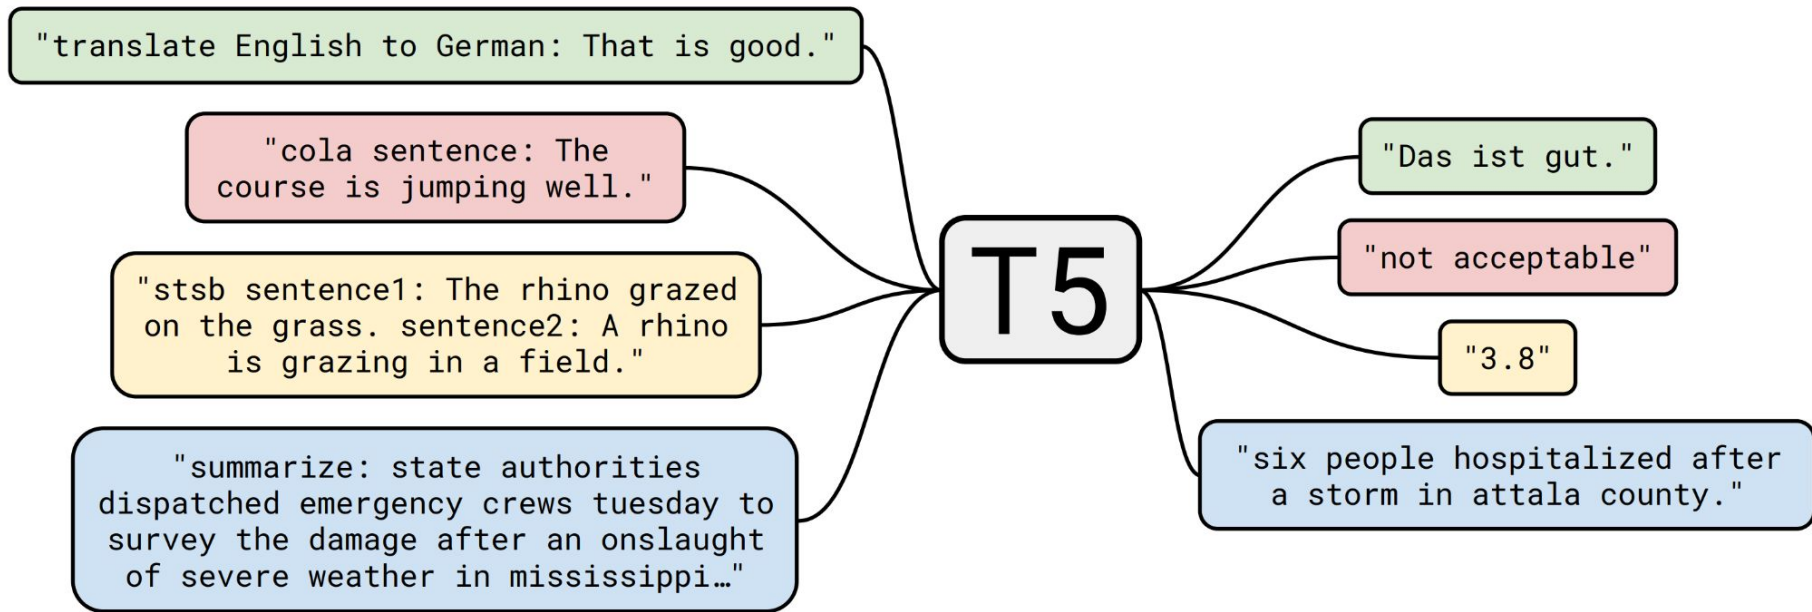


**Transformer decoder**  
(masked)


<BOS> positive

*decoder*

# T5 facilitates multitask learning




# Decoder-only model



***the architecture  
used in frontier  
LLMs***


**Transformer decoder  
(masked)**




## Note on cross-attention

- Can be used to inject non-text data (e.g., images, structured data, or even sensor readings) into the model

**Increasing model size enhances performance and sample efficiency**




# ... and unlocks new capabilities



*From "PaLM: Scaling Language Modeling with Pathways" by Chowdhery et al. (2022)*

The trend has continued to push the boundaries of possibility in NLP






# Inverse scaling

- <https://www.lesswrong.com/posts/iznohbCPFkeB9kAJL/inverse-scaling-prize-round-1-winners>
- <https://www.lesswrong.com/posts/DARiTSTx5xDLQGrrz/inverse-scaling-prize-second-round-winners>

# Inverse scaling (cont'd)



*Repeat my sentences back to me.*

*Input: I like dogs.*

*Output: I like dogs.*

*Input: What is a potato, if not big?*

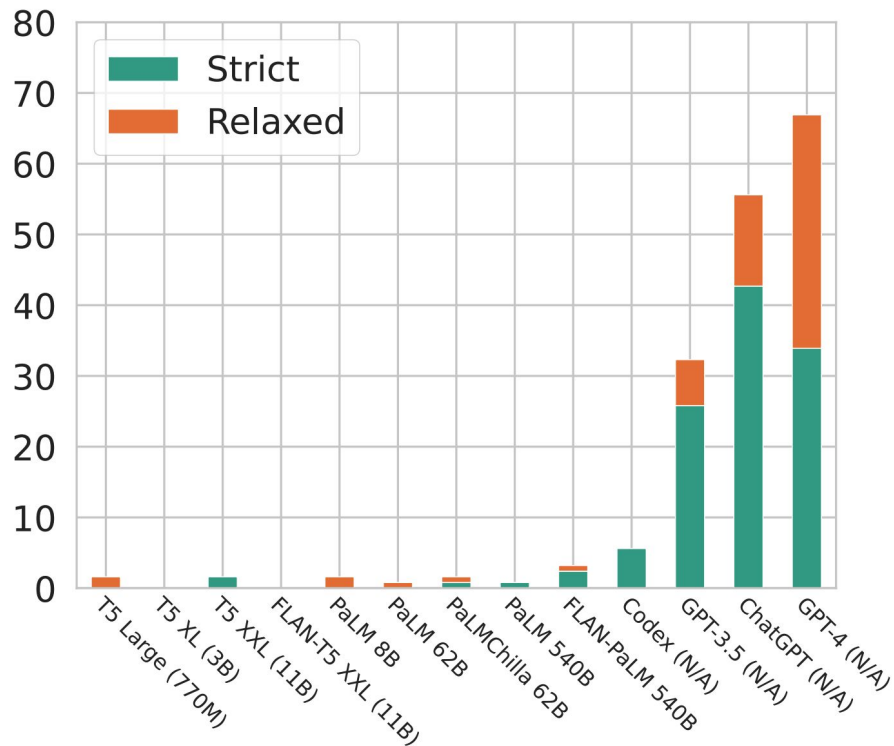
*Output: What is a potato, if not big?*

*Input: All the world's a stage, and all the men and women merely players. They have their exits and their entrances; And one man in his time plays many pango*

*Output: All the world's a stage, and all the men and women merely players. They have their exits and their entrances; And one man in his time plays many*

(where the model should choose 'pango' instead of completing the quotation with 'part.')

# False premise questions: When did Google release ChatGPT?



**False-premise questions**

Vu et al. 2023:  
<https://arxiv.org/abs/2310.03214>

# What can we scale?

- The loss scales as a power-law with:
  - **N**: model size
  - **D**: dataset size
  - the amount of compute used for training (e.g., number of training steps)

$$\text{Total Steps} = \frac{\text{Dataset Size} \times \text{Epochs}}{\text{Batch Size}}$$

Where:

- **Dataset Size:** Total number of training examples.
- **Epochs:** Number of times the model sees the entire dataset.
- **Batch Size:** Number of samples per batch update.

## Given a fixed compute budget, what is the optimal model size and dataset size for training?

Let's say you can use one GPU for one day

- Would you train a 5 million parameter LM on 100 books?
- What about a 500 million parameter LM on one book?
- Or a 100k parameter LM on 5k books?


Given a fixed compute budget, what is the optimal model size and dataset size for training?

- Kaplan et al. 2020
- **Chinchilla** (Hoffmann et al. 2022)



## Observations from Kaplan et al., 2020

- Performance depends largely on scale (model size, data size, and compute) and weakly on model architecture (e.g., depth, width)
- Performance improves most when model and dataset size scale together; increasing one while keeping the other fixed results in diminishing returns
-




## Issues with Kaplan laws

- Used same learning rate schedule for all training runs, regardless of how many training tokens / batches!
- This schedule needs to be adjusted based on the number of training steps; otherwise, it can impair performance
- The resulting “scaling laws” from Kaplan et al., are flawed because of this!



# Chinchilla scaling laws

- **Kaplan et al., 2020: prioritize increasing model size over data size**
  - With a 10x compute increase, increase model size by 5x and data size by 2x
  - With a 100x compute increase, model size 25x and data 4x
- **Chinchilla (Hoffmann et al., 2022): increase model and data size at the same rate**
  - With a 10x compute increase, increase both model size and data size by 3.1x
  - With a 100x compute increase, both model and data size 10x


For a given FLOP budget there is an optimal model to train



# Projecting optimal model size and number of tokens for larger models



# Large models should be significantly smaller and trained for much longer than is currently done (2022)



# Large models should be significantly smaller and trained for much longer than is currently done (2022)

Model	Size (# Parameters)	Training Tokens
LaMDA (Thoppilan et al., 2022)	137 Billion	168 Billion
GPT-3 (Brown et al., 2020)	175 Billion	300 Billion
Jurassic (Lieber et al., 2021)	178 Billion	300 Billion
Gopher (Rae et al., 2021)	280 Billion	300 Billion
MT-NLG 530B (Smith et al., 2022)	530 Billion	270 Billion
<i>Chinchilla</i>	70 Billion	1.4 Trillion

- $N$  – the number of model parameters, *excluding all vocabulary and positional embeddings*
- $C \approx 6NBS$  – an estimate of the total non-embedding training compute, where  $B$  is the batch size, and  $S$  is the number of training steps (ie parameter updates). We quote numerical values in PF-days, where one PF-day =  $10^{15} \times 24 \times 3600 = 8.64 \times 10^{19}$  floating point operations.

PF: PetaFLOP



# Gopher vs. Chinchilla

Random	25.0%
Average human rater	34.5%
GPT-3 5-shot	43.9%
<i>Gopher</i> 5-shot	60.0%
<b><i>Chinchilla</i> 5-shot</b>	<b>67.6%</b>
Average human expert performance	89.8%
June 2022 Forecast	57.1%
June 2023 Forecast	63.4%

Table 6 | **Massive Multitask Language Understanding (MMLU)**. We report the average 5-shot accuracy over 57 tasks with model and human accuracy comparisons taken from [Hendrycks et al. \(2020\)](#). We also include the average prediction for state of the art accuracy in June 2022/2023 made by 73 competitive human forecasters in [Steinhardt \(2021\)](#).

## Chinchilla's loss function

$$\hat{L}(N, D) \triangleq E + \frac{A}{N^\alpha} + \frac{B}{D^\beta}.$$

**Thank you!**